

Analytics of Realtime Soccer Match Sensor Data with JavaScript and WebGL—Reprocessed and Visualized for Web Browser or Command Line Consumption

Martin Kleppe
Ubilabs GmbH
Juliusstr. 25
22769 Hamburg, Germany
kleppe@ubilabs.net

ABSTRACT

In this paper, we report on a Web application with an additional command line interface capable of providing complex analyses over high velocity soccer match sensor data that was implemented in JavaScript and the Web Graphics Library (WebGL). This application visualizes analysis results graphically in the Web browser and in parallel also streams aggregated statistics in realtime to a command line interface. The data analyzed in this paper consists of raw sensor data that was recorded during an actual soccer match using wireless sensors embedded in the ball and the players' shoes. The thereof generated realtime analyses are twofold: on the one hand, they focus on the continuous computation of statistics such as ball possession, shots on goal, or player of the match, which is relevant to passive spectators like fans in a stadium or TV viewers at home. On the other hand, they focus on active observers of the match like team coaches or team managers that require more detailed analyses like running path visualizations, position heatmaps, *etc.*

Keywords

Realtime Analyses, Sensor Data, Data Streams, JavaScript, WebGL, Command Line Interface, Sports, Soccer

1. INTRODUCTION

Detailed sports game analyses are of high interest and relevance in today's professional sports leagues. Spectators are provided with additional statistics such as the number of shots on the goal, movement analyses, or percentage of ball possession per player or team. Furthermore, more detailed statistics provide useful information for coaches and team managers about players' performance during the match or in certain situations and also give insights about opponents, which could lead to modification in tactics. Although automated solutions such as high resolution video analyses

are desirable as they generate the required detailed statistics quickly, at present, most sports game statistics are still processed manually. Unfortunately, insights gained through image-based solutions are limited by image resolution, frame rate, and last not least prohibitive costs. For the ACM DEBS 2013 Grand Challenge,¹ the Fraunhofer Institute for Integrated Circuits (IIS)² have set up a realtime locating system on a soccer field in a stadium in Nuremberg, Germany. Every player and the ball were equipped with wireless sensors that produce high velocity sensor data at a total rate of about 15,000 position events per second.

In the following, we will outline our submission to the challenge that uses continuous computation of statistics in JavaScript to generate interactive visualizations and realtime analyses of the game such as ball possession, shots on goal, running analyses of all players and the two teams. The chosen approach can be called *hybrid*, as the same system is used to visualize the game in a Web browser³ and to output multiple event streams on the file system.

A short screencast of our application is available online at <http://bit.ly/debs-challenge-video> and the source code of our challenge contribution can be found on GitHub at the URL <https://github.com/ubilabs/soccer-debs-challenge/>.

The remainder of this paper is structured as follows. We cover related work in Section 2. In Section 3, we outline the methodology we used to approach the challenge problems. Section 4 is dedicated to implementation details and performance optimizations. We conclude with an outlook on future work in Section 5.

2. RELATED WORK

Automated sport analyses heavily depend on video systems that capture the game, compute differences between images and then use the remaining color information to track players and ball [8, 9]. Another approach is to equip the players with sensors that collect position data over time. This data can be combined with video processing, which can serve, for example, to select and zoom in on a situation where a certain player is within the opponent's penalty area [10]. Additionally, biometric sensors that collect information about the players' conditions—such as heartbeats

¹DEBS Challenge: <http://www.orgs.ttu.edu/debs2013/>

²Fraunhofer IIS <http://www.iis.fraunhofer.de/>

³We have tested our application in Google Chrome, version 26 for Mac OS X: <https://www.google.com/chrome>

and body temperatures—provide data that is used to analyze the performance during the game [1]. Spatial game analytics (*e.g.*, heatmaps that display the players’ distribution over time) are widely recognized by sports professionals as one of the most useful applications, because they can be used to optimize the team distribution for a specific game [6].

Using JavaScript to analyze enormous amounts of data in realtime—as it is the case with high accuracy sensory data—was not possible until Node.js,⁴ a server-side software system designed for writing scalable Internet applications, notably Web servers. Node.js programs are written on the server side in JavaScript, using event-driven, asynchronous I/O to minimize overhead and maximize scalability. Node.js is a packaged compilation of Google’s V8 JavaScript engine⁵ and was shown to be suitable for building high-performance network programs [14] thanks to its event-driven and non-blocking nature. Furthermore, HTML5 features such as Canvas and WebSockets are used for realtime monitoring systems [3].

When visualizing three-dimensional content, WebGL [11],⁶ a Web standard for a low-level 3D graphics API based on OpenGL ES 2.0,⁷ is the tool of choice to render 3D objects in the browser. It is used for MMOGs (Massively Multiplayer Online Games) [4], efficient rendering of 3D models [13], and interactive visualization of volumetric data [5].

3. PROBLEM APPROACH

Our present submission to the ACM DEBS 2013 Grand Challenge started with the question: “is it possible to read and analyze the provided input data stream and visualize it in the browser?” In a first step, initial tests turned out that parsing the provided file with the raw sensor data and simply displaying the positions of all players and balls in 3D is possible with a factor of twenty times of the actual speed on a standard consumer laptop, *i.e.*, a minute of the real game is replayed within just three seconds. This leaves enough time to make additional calculations. The two teams and the ball can easily be rendered recognizable by using different colors for each. Additionally, the soccer field and goals are drawn to visually check different game situations.

In a second step, we implemented the required queries. For every piece of information, a visual element was added to the browser interface to keep track of the computations and avoid gross errors. These visual elements include: a tracing line to highlight the direction of the ball movement, the list of players with statistics about ball possession, colored sparklines for running analyses, the precalculated ball path, the potential hit target, highlighting of the current player, an animated acceleration bar, and the current time. A screenshot of the application can be seen in Figure 1.

In a third step, as the application has to deal with large datasets, it is critical to observe the internally used memory. The Chrome Developer Tools Timeline panel⁸ was used to detect if any of the involved scripts result in increasing memory usage without freeing it at the script’s end. If such memory leaks were detected, they were carefully evaluated and the memory-leaking code improved.

⁴Node.js: <http://nodejs.org/>

⁵V8: <https://code.google.com/p/v8/>

⁶WebGL: <http://webgl.org/>

⁷OpenGL: http://www.khronos.org/opengles/2_X/

⁸Chrome Developer Tools: <https://developers.google.com/chrome-developer-tools/docs/timeline>

In a final step, after all visual parts were included, the original code was extended to also run via the command line without the need of a Web browser. Therefore, a bridge was created to either render HTML or output several file streams if executed on the command line. The result is an application that does both, visualize the match in a browser or store aggregated results on the hard disk.

4. IMPLEMENTATION DETAILS

Input Data Format.

The original data stream was captured during the game and resulted in a 4.62 GB CSV input file. Position updates for sensors in players’ shoes and goalkeepers’ hands are provided with a frequency of 200Hz. The sensors in all balls update with 2000Hz. A sensor record contains the following data: sensor id, timestamp in picoseconds, position (in a three-dimensional coordinate system) of the sensor in millimeters, $|v|$ (in $\mu\text{m/s}$), v_x , v_y , v_z describe the speed of the direction of objects as a vector, and $|a|$ (in $\mu\text{m/s}^2$), a_x , a_y , a_z describe the absolute acceleration and its constituents in three dimensions.

Implemented Queries.

Based on this data, the following queries are required for the ACM DEBS 2013 Grand Challenge: running analysis, ball possession, heatmaps, and shots on goal. Furthermore, goal detection and detection of when the ball is out were implemented as additional features of the application.

Data Flow.

All entries of the original data stream are distributed to several JavaScript objects based on a mapping table that includes more information about the sensor type. Whenever a ball position update was detected, the system uses the last known position to check for a goal or whether the ball has left the field. If the ball acceleration peaks, it detects the associated player and computes the shot target based on the current speed vector and gravity. Position updates for all players are collected for the purpose of running statistics and heatmap calculations.

Performance Tuning.

To optimize performance and to avoid long running scripts, plain JavaScript with almost zero dependencies was used, with the exception of the already highly performance-optimized libraries fishbone.js⁹ and Three.js;¹⁰ the former being an extremely lightweight JavaScript library with automatic method chaining, automatic context binding, event support, and simple inheritance and the latter being a wrapper for low-level WebGL instructions. The code was organized in two kinds of modules: (i) simple class-like modules with prototype-based inheritance that were used for fast-changing game objects like players or ball on the one hand and (ii) modules that handle events between these objects and the streams on the other. A centralized runner script handles most of the time-consuming calculations within a flat lexical scope to avoid nested functions and variable lookups. This was especially critical for large loops that occur when the input stream emits new records.

⁹Fishbone.js: <https://github.com/aemkei/fishbone.js>

¹⁰Three.js: <http://threejs.org/>

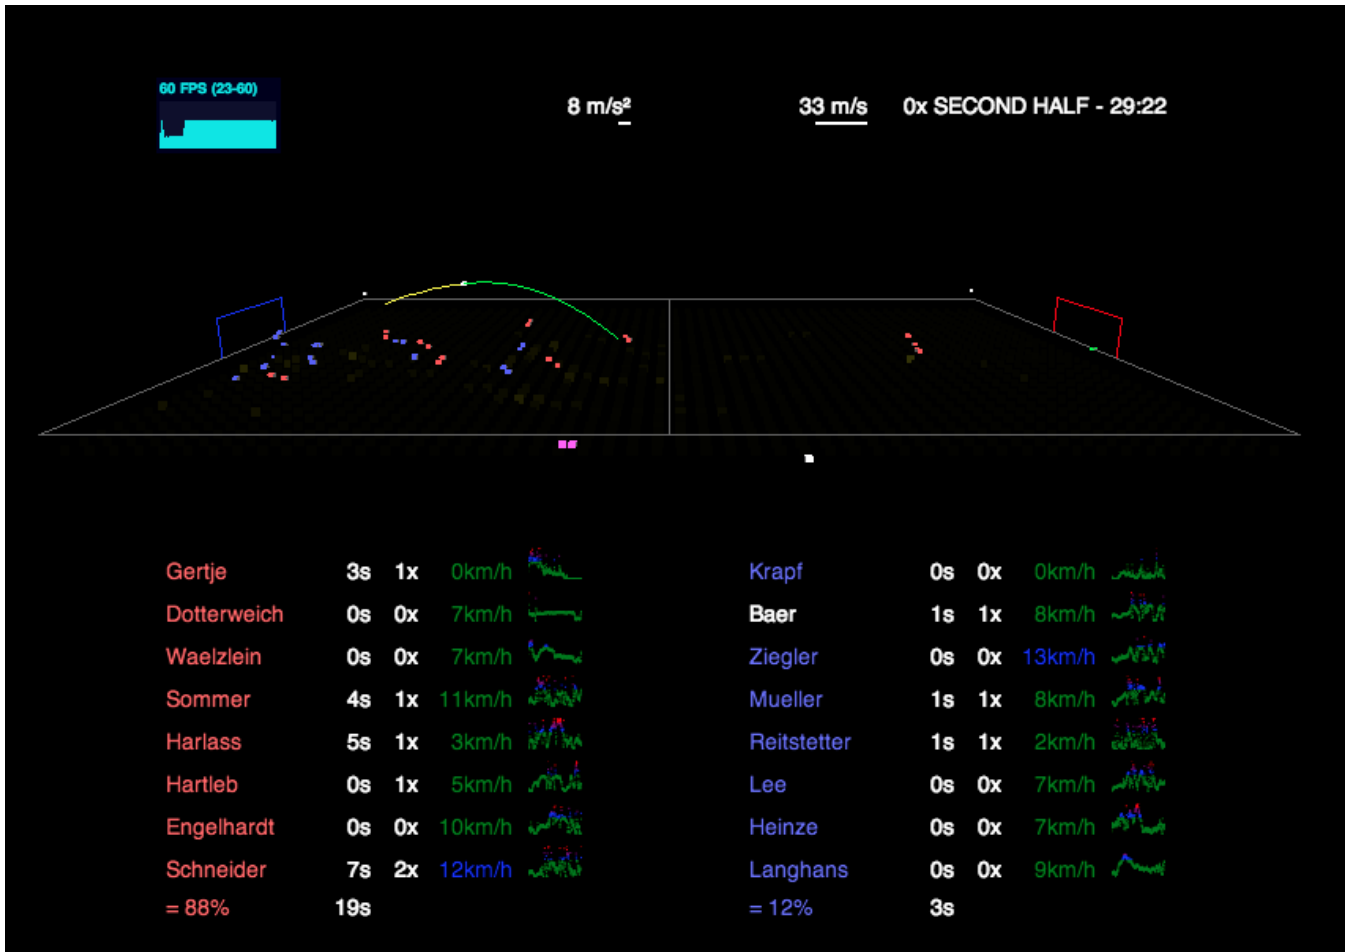


Figure 1: Screenshot of the application with a tracing line to highlight the direction of the ball movement (yellow), the precalculated ball path (green), the list of players with individual statistics about ball possession, colored sparklines for running analyses, the potential hit target, current player highlighting, animated acceleration bar, and current time

Data Aggregation.

Whenever an update of one of the (several during the match utilized) balls is detected, the program evaluates if the current ball also is the currently active match ball by comparing its position with the field's boundaries. Then the nearest player is selected and if the ball's acceleration peaks, shots on goal and ball possessions (per player and team) are evaluated. With a frequency of 50Hz, the player's current position is recorded into a large array to generate running statistics and heatmaps based on different time frames (1, 5, 20 minutes and the whole game). This is done by looping through the records multiple times per interval and comparing all dimensions to create aggregated values.

Visualization.

To visualize the results in the browser, position properties of JavaScript objects are updated whenever new data arrives. As WebGL is a state machine [2], these updates are handled very fast. Geometries for sensors are displayed as colored cubes, the field and ball paths are simple polylines and the heatmap is a particle system. We use rectangular sprites, as the rendering performance of 2D sprites is considerably better than updating complex geometries [7]. The list of players is drawn as an unordered HTML list (``)

and colors are assigned via CSS.¹¹ The color-coded sparkline graphs at the end of each list item are drawn using HTML5 Canvas,¹² which was shown to perform better than plain HTML, SVG, and WebGL, as it generates highly optimized stateless bitmaps.

Command Line Output.

Output for the command-line version was implemented using so-called file streams: for every calculation that emits events, a writable stream¹³ is created in Node.js. The current implementation pipes the output to multiple files on the hard disk, each for every query type: Player running analysis (1 stream), aggregated running statistics (4 time frames), player ball possession (1 stream), team ball possession (1 stream), heatmaps (4 time frames), and shots on goals (1 stream). This results in a total of 12 files written to disk. The complete calculation of the entire match data takes about 600 seconds on our testing device.

¹¹CSS: <http://www.w3.org/Style/CSS/>

¹²HTML5 Canvas: <https://dev.w3.org/html5/2dcontext>

¹³Node.js Writable Stream: http://nodejs.org/api/all.html#all_class_stream_writable

5. CONCLUSIONS AND FUTURE WORK

In this paper, we have shown that modern JavaScript engines such as V8 are well suited to process large amounts of data in realtime. With the current two-fold implementation, it is possible to analyze a full soccer match on the command line and to visualize it in the browser. HTML5 features such as Canvas and WebGL draw graphics without any noticeable performance degradation, and the event-driven, non-blocking I/O model of Node.js was shown to be an efficient way to read, process, and write data. As mentioned above, the current implementation outputs all streams to the file system. However, using the abstract stream pattern in Node.js, they can be piped [12] to other types of streams such as WebSockets.¹⁴

In the future, this will enable the new use case of allowing for mobile devices such as cell phones and tablets with limited storage and computing power to be used for the visualization. They connect to the main server via a WebSocket connection and receive only small chunks of data that are necessary to render relevant information. Coaches and team members can benefit from such a solution for mobile devices, as they are small and portable compared to laptop computers. Further future work can include to port the current version of the application to other ball team sports, *e.g.*, American football, rugby or basketball. More automatic analyses can be integrated, such as number of corner shots, passes, and duel statistics. If combined with additional biometric sensor data (pulse, power) this could give insights into the current player's condition or even training progress during a whole season.

Acknowledgments

We would like to thank the following persons for their valuable comments: Thomas Steiner, Nenad Nikolić, Klaus Trainer, Michael Pletziger, Jens Wille, Samuel Oey, Robert Katzki, Astrid Schildkopf.

6. REFERENCES

- [1] R. A. Alonso. Sports Telemetry System for Collecting Performance Metrics and Data, 2010. Patent US20100283630.
- [2] D. Cantor. *WebGL Beginner's Guide*. Packt Pub Limited, 2012.
- [3] S.-H. Cha. Developing a HTML5-Based Real-Time Monitoring System for Wireless Sensor Networks. In *Advances in Mechanical and Electronic Engineering*, pages 241–245. Springer, 2013.
- [4] G. Chen. WebGL Massively Multiplayer Online Game.
- [5] J. Congote, A. Segura, L. Kabongo, A. Moreno, J. Posada, and O. Ruiz. Interactive Visualization of Volumetric Data with WebGL in Real-time. In *Proceedings of the 16th International Conference on 3D Web Technology*, pages 137–146. ACM, 2011.
- [6] A. Drachen and M. Schubert. Spatial Game Analytics. In M. Seif El-Nasr, A. Drachen, and A. Canossa, editors, *Game Analytics*, pages 365–402. Springer London, 2013.
- [7] R. C. Hoetzlein. Graphics Performance in Rich Internet Applications. *Computer Graphics and Applications, IEEE*, 32(5):98–104, 2012.
- [8] Y. Huang, J. Llach, and S. Bhagavathy. Players and Ball Detection in Soccer Videos based on Color Segmentation and Shape Analysis. In *Proceedings of the 2007 International Conference on Multimedia Content Analysis and Mining, MCAM'07*, pages 416–425, Berlin, Heidelberg, 2007. Springer-Verlag.
- [9] D. Liang, Y. Liu, Q. Huang, and W. Gao. A Scheme for Ball Detection and Tracking in Broadcast Soccer Video. In *Proceedings of the 6th Pacific-Rim Conference on Advances in Multimedia Information Processing – Volume Part I, PCM'05*, pages 864–875, Berlin, Heidelberg, 2005. Springer-Verlag.
- [10] N. R. Løvhaug and E. Lie. Fully Automated Football Analysis, 2012. http://www.forskningsradet.no/en/Newsarticle/Fully_automated_football_analysis/1253981215834/p1177315753918, accessed 04/21/2013.
- [11] C. Marrin. WebGL Specification, Version 1.0.1, 27 January 2012, 2012. <https://www.khronos.org/registry/webgl/specs/1.0/>, accessed 04/21/2013.
- [12] M. Ogden. Node Streams: How do they work?, 2012. <http://maxogden.com/node-streams>, accessed 04/21/2013.
- [13] B. Sawicki and B. Chaber. Efficient Visualization of 3D Models by Web Browser. *Computing*, pages 1–13, 2013.
- [14] S. Tilkov and S. Vinoski. Node.js: Using JavaScript to Build High-Performance Network Programs. *Internet Computing, IEEE*, 14(6):80–83, 2010.

¹⁴WebSockets: <http://dev.w3.org/html5/websockets/>